

# Distributed Systems: Final Project

## TLDR

The project for this course will be to develop, a distributed, reliable backend in support of a mildly complex (think twitter) social media application.

The development of this project is divided into 2 stages! The stages are due **November 29th, and December 13th**, respectively. This project can be done individually or in pairs. If you work in pairs, we will check the git history to see contribution is roughly 50/50!!

Your project will be split into 2 stages, the stages are as follows:

1. A web application split into several services, with a *stateless* web server and *at least one* backend service, which persist data and communicate with the stateless web server via gRPC. The web application exposes a simple social media website.
2. The backend service(s) are now stateless, persisting their state in a raft, replicated data store.

## Important Details

### Anubis LMS (required):

For this assignment we will be using the new Anubis Learning Management System. Anubis will generate a github repository for you from a template for the lab. You can then submit your homework by pushing to the repo. All push events are timestamped by the Anubis system.

This system will also provide you with optional Cloud IDEs servers at a click of a button for you to work on your homework on. These Cloud IDE servers will have all the necessary software for you to be able to complete this lab. The Cloud IDE servers have an “autosave” feature where a bot user will push any unsaved changes to github every 5 minutes. You can also type the autosave command in the Cloud IDE terminal to manually trigger an autosave. You are responsible for making sure that your work ends up submitted on github on time. Please take a minute to understand and double check that your work is getting pushed to github on time.

If you do not wish to use the Cloud IDEs, you can clone the repo that Anubis generates for you and push to it from your local machine. You can submit (by pushing to your repo) as much as you would like up to the assignment deadline.

To use Anubis: 1. Head to <https://anubis.osiris.services/> 1. Press “LogIn” on the side bar, and log in with your NYU ID 1. Enter GitHub username on the profile page (If not done yet) 1. Go to “Courses”, and press “Join Class” with code ilovedistsys to register for the class 1. Go to “Assignments” - Press “Create

Assignment Repo” (This will create a repository on Github) - Start and use an Anubis Cloud IDE, or clone the generated repo locally. Please make sure that you are pushing to github (and therefore submitting) before the deadline.

## Stage 1 - (November 29)

### Deliverable

In this stage you will start by creating a simple web application, comprised of a web server written in Go, serving html files from a single machine. Once you have that working, you will split off your backend into a separate service. It can be a monolith (i.e. one service that performs all of the tasks related to your system’s state management), or a set of services that each perform one small task, as described in-class.

This service (set) must communicate with the web server using an RPC Framework like gRPC (strongly recommended!) or Thrift. If you run into any issues getting these frameworks set up, don’t hesitate to contact the TAs!

At this point, your webpage service (i.e. the service that is rendering templates / serving web pages) should be *stateless*. In practical terms, you should be able to horizontally scale your web service without causing any concurrency problems on the backend, or UX problems on the front end. :wave: **The Web Service should no longer be persisting any data (even session data!!)**, but instead be fully reliant on the backend service. If this isn’t the case, you still have work to do :)

### Features

This application needs to have a small number of very clear functions. The minimum set are as follows:

1. Creating an account, with username and password
2. Logging in as a given user, given username and password
3. Users can follow other users. This action must be reversible.
4. Users can create posts that are associated with their identity.
5. Users can view some feed composed only of content generated by users they follow.

Note that these operations *should* be accomplished efficiently, but *need* not be. You won’t lose points for  $O(n)$  algorithms, where an  $O(\log n)$  solution is possible. However, you’re selling yourself short ;)

If you want to build a system other than this twitter example, that is totally OK. Just speak to the professor/TA first, to get approval. We just want to make sure that your application is of comparable difficulty and testability.

### Tests

You *will* be expected to write tests for your functions!! If you encounter any issues with testing, don't hesitate to ping your TA! If you want to use fancy testing frameworks (like [ginkgo](#)), that's awesome! Not required, but might be fun.

To be clear we're not setting an arbitrary code coverage metric, but tests are super important, so they'll factor into your team's grade. If you're nervous about whether your tests are good enough, speak to your TA about it. We just want to drive home that tests are super important, especially in this phase, as these tests will be your canaries when you add your a persistent backend in step 2.

## Frameworks and Vendoring

On that note, if you want to use frameworks such as buffalo or ginkgo, go ahead, but make sure you practice [vendoring them as dependencies](#)! It'll make it easier on your teammates, and easier for me to grade! That being said, you definitely will not *need* to use any of these frameworks. As our system is relatively small, they won't save you a ton of time. It's also worth learning how std packages like `net/http` and `testing` work, as that's what all of these frameworks are built on.

## Structure

As a friendly advice, you should try to structure your application as follows. You don't *need* to, but it will definitely make stage 2 easier! we have created some empty web.go files in the template to encourage you to use this template.

```
cmd
|-- web
|   `-- web.go --> build target (the code that actually runs the server)
web
|-- web.go
|-- config.go --> config for web, determines port to bind to, etc.
|-- cmd      --> implementation of build target
`-- auth
    |-- auth.go --> authentication module, creates a new auth object,
    |             starts it, stops it.
    |-- config.go --> token validity duration, hash difficulties, etc.
    |-- errors.go --> errors that auth object can reply with
    |-- *.go      --> implement auth; use contexts, generated
    |             proto, and storage packages
    |-- *_test.go --> don't forget tests!
    |-- storage
    |   |-- storage.go --> storage interface for auth
    |   `-- memory
    |       |-- config.go
    |       `-- memory.go --> *threadsafe* implementation,
```

```

                                using maps and lists
    `-- authpb
       `-- models.proto --> used to generate all data
                                primitives used by auth module

```

The auth package is enumerated as an example. All of auth's sister packages should look pretty similar, and **web** should compose these packages, hiding access to them behind http handlers, as was shown in class.

If you want, you can totally abandon this structure and follow your own path!! It's totally up to you, but please pick something sensible!! :)

## Resources

The following resources might be helpful:

1. This [open source book](#), which explains how to build a go web server from fundamentals. Note that there is way more information in this link than you'll need, but it's a great overview. Some of the snippets in Chapter 3 might be especially useful.
2. This [blog post](#) explains a popular method of organizing projects.
3. This [exploration of Kubernetes source](#), which is really cool, and gives you an intimate sense of how the project works.
4. This [example travis script](#) for those of you who want to play around with automated builds and tests using Travis!
5. For leveraging contexts, check out [this blog post](#), which goes over the subject way better than I could!!
6. This might be a bit late to slide in, but [this talk](#) on concurrency patterns is great.
7. For faking http responses, [this package](#) is great.
8. [This document](#) is a super concise explanation of installing gRPC and protobufs (the transport serialization used by gRPC). It also references an example, which is a great resource in and of itself for learning protobuf syntax. In particular, it contains snippets for creating, registering, and dialling with gRPC servers/clients, which you can directly use in the project.
9. Adam's [gen proto script](#) :)
10. Your TA! If you have any issues!

## Stage 2 - (December 13)

### Summary

In stage 2, we are finally going to give our system a much-needed upgrade – persistence!! And not just any kind of persistence – we're talking about highly-available, replicated persistence!

To accomplish this, we're going to use one of the two popular open source Raft implementations out there, either the one from [CoreOS](#) or [Hashicorp](#). You can

choose to use whichever one you prefer; in your final presentation, you will need to give some rationale for why you picked the one you did! :)

## Deliverables

This is a relatively small step, but it requires learning how your raft implementation of choice looks like!! It should end up being more of an ops problem than a dev problem. Again, if you've followed the example architecture, it shouldn't be a huge headache. If you find that it is, contact your TA!!

You'll need to:

1. Write code that stands up raft nodes (on your local machine is ok!!)
2. Replace your tried and tested storage implementations with a new one, which wraps a raft client for your raft implementation of choice. This will live alongside your old implementation in the **storage** directory, if you're following the recommended structure. This involves code that establishes a connection with your raft cluster. Now, you'll find that your storage implementation's **config**, **start** and **shutdown** actually do some heavy lifting!!
3. Submit all the code into github
4. Create a video presentation 5-10 mins with your teammate

## Details

Each service should have its *own client* wrapper for contacting the raft cluster, but it is totally OK if they are all targeting the same cluster. If you want, you can target multiple clusters – good luck getting all those to run on your laptop, though :).

Don't worry about hosting this solution anywhere other than your laptop. If you want to host it on AWS, heroku, or what have you (just to show off!), that's totally OK. Just make sure that you have a configuration that will also work on your local machine. One hack for accomplishing this in the recommended structure, is to build out **cmd** directory as follows:

```
cmd
|-- local
|   |-- web
|   `-- backend
`-- aws
    |-- web
    `-- backend
```

In the above example, your config initializations under **local** would look very different from your config inits under **prod**. Alternatively, you can read config objects values from environment variables, have different env-setting shell scripts for aws/local, and keep a single **cmd** folder. It's really up to you!

## **Academic Honesty**

Aside from the narrow exception for collaboration on homework, all work submitted in this course must be your own. Cheating and plagiarism will not be tolerated. If we suspect of plagiarism you will get a 0 on the assignment and potentially the class. If you have any questions about a specific case, please ask me. We will be checking for this! NYU's Policy on Academic Misconduct: <https://engineering.nyu.edu/student-life/student-activities/office-student-affairs/policies/code-conduct>

## **Errata**

Did you find some bug with this document? Want to improve it? Go ahead and submit an issue on [Github](#).